

# Part\_I\_exploration

September 3, 2022

## 1 Part I - Ford GoBike System Data Exploration

### 1.1 by Philip Obiorah

### 1.2 Introduction

Ford GoBike System Data: This data set includes information about individual rides made in a bike-sharing system covering the greater San Francisco Bay area. .

### 1.3 Preliminary Wrangling

```
[1]: # import all packages and set plots to be embedded inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
import datetime as dt
from datetime import datetime as dt
%matplotlib inline
```

```
[2]: #load the dataset into a pandas dataframe
ford_go_bike = pd.read_csv('201902-fordgobike-tripdata.csv')
ford_go_bike.sample(15)
```

```
[2]:
```

	duration_sec	start_time	end_time	\
12847	415	2019-02-27 17:12:36.9520	2019-02-27 17:19:32.8970	
181506	867	2019-02-01 09:01:34.1620	2019-02-01 09:16:01.8750	
90577	783	2019-02-16 12:04:13.0790	2019-02-16 12:17:16.2600	
55846	463	2019-02-21 09:32:56.0270	2019-02-21 09:40:39.3280	
94733	890	2019-02-15 15:48:17.0750	2019-02-15 16:03:08.0570	
95428	972	2019-02-15 13:39:34.9770	2019-02-15 13:55:47.0640	
62167	799	2019-02-20 17:04:06.0900	2019-02-20 17:17:25.2810	
129421	1164	2019-02-10 10:15:28.6390	2019-02-10 10:34:53.2290	
22742	690	2019-02-26 06:54:34.4100	2019-02-26 07:06:04.4960	
173242	294	2019-02-03 12:42:02.0430	2019-02-03 12:46:56.7190	
5644	227	2019-02-28 12:39:08.1920	2019-02-28 12:42:55.3520	

169510	269	2019-02-04 09:58:09.9640	2019-02-04 10:02:39.2100
108216	554	2019-02-13 09:11:21.5940	2019-02-13 09:20:36.3070
155609	158	2019-02-06 08:46:07.9810	2019-02-06 08:48:46.5630
161703	1094	2019-02-05 13:50:29.9470	2019-02-05 14:08:44.9430

	start_station_id	start_station_name \
12847	8.0	The Embarcadero at Vallejo St
181506	194.0	Lakeshore Ave at Trestle Glen Rd
90577	370.0	Jones St at Post St
55846	60.0	8th St at Ringold St
94733	36.0	Folsom St at 3rd St
95428	160.0	West Oakland BART Station
62167	240.0	Haste St at Telegraph Ave
129421	76.0	McCoppin St at Valencia St
22742	101.0	15th St at Potrero Ave
173242	27.0	Beale St at Harrison St
5644	241.0	Ashby BART Station
169510	202.0	Washington St at 8th St
108216	126.0	Esprit Park
155609	90.0	Townsend St at 7th St
161703	86.0	Market St at Dolores St

	start_station_latitude	start_station_longitude	end_station_id \
12847	37.799953	-122.398525	22.0
181506	37.811081	-122.243268	7.0
90577	37.787327	-122.413278	79.0
55846	37.774520	-122.409449	67.0
94733	37.783830	-122.398870	84.0
95428	37.805318	-122.294837	155.0
62167	37.866043	-122.258804	189.0
129421	37.771662	-122.422423	10.0
22742	37.767079	-122.407359	3.0
173242	37.788059	-122.391865	30.0
5644	37.852477	-122.270213	168.0
169510	37.800754	-122.274894	219.0
108216	37.761634	-122.390648	125.0
155609	37.771058	-122.402717	67.0
161703	37.769305	-122.426826	70.0

	end_station_name \
12847	Howard St at Beale St
181506	Frank H Ogawa Plaza
90577	7th St at Brannan St
55846	San Francisco Caltrain Station 2 (Townsend St...
94733	Duboce Park
95428	Emeryville Public Market
62167	Genoa St at 55th St

129421	Washington St at Kearny St
22742	Powell St BART Station (Market St at 4th St)
173242	San Francisco Caltrain (Townsend St at 4th St)
5644	Alcatraz Ave at Shattuck Ave
169510	Marston Campbell Park
108216	20th St at Bryant St
155609	San Francisco Caltrain Station 2 (Townsend St...
161703	Central Ave at Fell St

	end_station_latitude	end_station_longitude	bike_id	user_type	\
12847	37.789756	-122.394643	1566	Subscriber	
181506	37.804562	-122.271738	1298	Subscriber	
90577	37.773492	-122.403672	5246	Subscriber	
55846	37.776639	-122.395526	3263	Subscriber	
94733	37.769200	-122.433812	6522	Subscriber	
95428	37.840521	-122.293528	5089	Subscriber	
62167	37.839649	-122.271756	5251	Subscriber	
129421	37.795393	-122.404770	5039	Subscriber	
22742	37.786375	-122.404904	6311	Subscriber	
173242	37.776598	-122.395282	4557	Subscriber	
5644	37.849595	-122.265569	6145	Subscriber	
169510	37.809824	-122.280192	5165	Subscriber	
108216	37.759200	-122.409851	4974	Subscriber	
155609	37.776639	-122.395526	4986	Subscriber	
161703	37.773311	-122.444293	4931	Subscriber	

	member_birth_year	member_gender	bike_share_for_all_trip
12847	1987.0	Female	No
181506	1974.0	Male	No
90577	1997.0	Female	No
55846	1989.0	Male	No
94733	1981.0	Female	No
95428	1978.0	Male	No
62167	1989.0	Male	No
129421	1991.0	Male	No
22742	1993.0	Male	No
173242	1989.0	Male	No
5644	1965.0	Female	No
169510	1993.0	Male	Yes
108216	1990.0	Male	No
155609	1988.0	Female	No
161703	1990.0	Other	No

## Issues

- Incorrect datatype for start\_time, end\_time, member\_birth, start\_station\_id, end\_station\_id, member\_birth\_year

- Missing values in member\_gender, member\_birth\_year, end\_station\_id, start\_station\_name, start\_station\_id and end\_station\_name
- No special column for age, day, and month
- Latitude and longitude are great features but we cannot immediately get insight for them without performing some calculations of the distance covered.

```
[3]: ford_go_bike.describe()
```

```
[3]:
```

	duration_sec	start_station_id	start_station_latitude \
count	183412.000000	183215.000000	183412.000000
mean	726.078435	138.590427	37.771223
std	1794.389780	111.778864	0.099581
min	61.000000	3.000000	37.317298
25%	325.000000	47.000000	37.770083
50%	514.000000	104.000000	37.780760
75%	796.000000	239.000000	37.797280
max	85444.000000	398.000000	37.880222

	start_station_longitude	end_station_id	end_station_latitude \
count	183412.000000	183215.000000	183412.000000
mean	-122.352664	136.249123	37.771427
std	0.117097	111.515131	0.099490
min	-122.453704	3.000000	37.317298
25%	-122.412408	44.000000	37.770407
50%	-122.398285	100.000000	37.781010
75%	-122.286533	235.000000	37.797320
max	-121.874119	398.000000	37.880222

	end_station_longitude	bike_id	member_birth_year
count	183412.000000	183412.000000	175147.000000
mean	-122.352250	4472.906375	1984.806437
std	0.116673	1664.383394	10.116689
min	-122.453704	11.000000	1878.000000
25%	-122.411726	3777.000000	1980.000000
50%	-122.398279	4958.000000	1987.000000
75%	-122.288045	5502.000000	1992.000000
max	-121.874119	6645.000000	2001.000000

```
[4]: #change start_time , end_time and member_birth_year datatypes to datetime
ford_go_bike['start_time'] = pd.to_datetime(ford_go_bike['start_time'])
ford_go_bike['end_time'] = pd.to_datetime(ford_go_bike['end_time'])
ford_go_bike['member_birth_year'] = pd.
↳to_datetime(ford_go_bike['member_birth_year'])
```

```
[5]: #test change start_time , end_time and member_birth_year datatypes to datetime
ford_go_bike.dtypes
```

```
[5]: duration_sec          int64
     start_time            datetime64[ns]
     end_time              datetime64[ns]
     start_station_id      float64
     start_station_name     object
     start_station_latitude float64
     start_station_longitude float64
     end_station_id        float64
     end_station_name       object
     end_station_latitude   float64
     end_station_longitude  float64
     bike_id               int64
     user_type              object
     member_birth_year      datetime64[ns]
     member_gender          object
     bike_share_for_all_trip object
     dtype: object
```

```
[6]: # Let examine duration sec to find out if any duration is more than 24 hrs
     ↪ (86400)
     ford_go_bike['duration_sec'].max()
```

```
[6]: 85444
```

Maximum duration is 85444 , so it is safe to assume that all rides occurred in one day. With that assumption we shall extract the day and month for each ride. Given the dataset we also assume that all rides occur in the same year 2019.

```
[7]: ford_go_bike['day'] = ford_go_bike['start_time'].dt.day_name()
     ford_go_bike['month'] = ford_go_bike['start_time'].dt.month_name()
```

```
[8]: # Converting float to int
     ford_go_bike.member_birth_year = ford_go_bike.member_birth_year.astype(np.int)
```

/tmp/ipykernel\_16761/1718661354.py:2: DeprecationWarning: `np.int` is a deprecated alias for the builtin `int`. To silence this warning, use `int` by itself. Doing this will not modify any behavior and is safe. When replacing `np.int`, you may wish to use e.g. `np.int64` or `np.int32` to specify the precision. If you wish to review your current use, check the release note link for additional information.

Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
     ford_go_bike.member_birth_year = ford_go_bike.member_birth_year.astype(np.int)
```

```
[9]: #Let define a function to allow us extract age form member_birth_year
def get_age(dob):
    today = dt.now().year
    return today - dob
```

```
[10]: # We use the get_age function that returns age to create a new column
      ↪ `member_age` which would hold member ages.
ford_go_bike["member_age"] = get_age(ford_go_bike["member_birth_year"])
```

```
[11]: #Let us examine the memeber_age column
ford_go_bike["member_age"].describe()
```

```
[11]: count      1.834120e+05
      mean      -4.156280e+17
      std       1.913313e+18
      min      -9.223372e+18
      25%       2.900000e+01
      50%       3.400000e+01
      75%       4.100000e+01
      max       1.440000e+02
      Name: member_age, dtype: float64
```

```
[12]: ford_go_bike[ford_go_bike["member_age"] > 70].count()
```

```
[12]: duration_sec      1032
      start_time        1032
      end_time          1032
      start_station_id  1032
      start_station_name 1032
      start_station_latitude 1032
      start_station_longitude 1032
      end_station_id    1032
      end_station_name  1032
      end_station_latitude 1032
      end_station_longitude 1032
      bike_id           1032
      user_type         1032
      member_birth_year  1032
      member_gender     1032
      bike_share_for_all_trip 1032
      day               1032
      month             1032
      member_age        1032
      dtype: int64
```

```
[13]: # Lets look at the 99% of bike riders age
ford_go_bike.member_age.describe(percentiles=[.99])
```

```
[13]: count    1.834120e+05
      mean    -4.156280e+17
      std     1.913313e+18
      min    -9.223372e+18
      50%     3.400000e+01
      99%     6.600000e+01
      max     1.440000e+02
      Name: member_age, dtype: float64
```

- It seems there is an outlier in the `member_age`. The maximum age is 144 which appears to be abnormal, and over 99% for all ages are below 70.
- With the above in view it would be safe to limit our analysis to ages below 70.

```
[14]: #select only records with member_age <= 70
      ford_go_bike = ford_go_bike.query("member_age <= 70")
```

```
[15]: #test member_age values are less than 70
      ford_go_bike.member_age.describe()
```

```
[15]: count    1.823800e+05
      mean    -4.179799e+17
      std     1.918462e+18
      min    -9.223372e+18
      25%     2.900000e+01
      50%     3.400000e+01
      75%     4.100000e+01
      max     7.000000e+01
      Name: member_age, dtype: float64
```

Calculate distance covered from Latitude and longitude Reference:  
<https://stackoverflow.com/questions/19412462/getting-distance-between-two-points-based-on-latitude-longitude>

```
[16]: # We would use # Will be using the math radians, sin, cos, sqrt and atan2
      from math import radians, sin, cos, sqrt, atan2
      # A function that would accept longitude stat and lat stat and long. end and
      ↪ lat. end
      # to return the total distance between two points.
      def distance(start, end):

          lat1, long1 = start
          lat2, long2 = end
          radius = 6371 # Approximate radius of the earth

          dlat = radians(lat2 - lat1)
          dlong = radians(long2 - long1)
```

```

a = (sin(dlat / 2) * sin(dlat / 2) + cos(radians(lat1)) *
↳cos(radians(lat2)) * sin(dlong / 2) * sin(dlong / 2))
c = 2 * atan2(sqrt(a), sqrt(1 - a))
distance = radius * c

# Round up the values to two decimal places
return round(distance, 2)

```

```

[17]: ford_go_bike['distance_km'] = ford_go_bike.apply(lambda x:
↳distance((x['start_station_latitude'], x['start_station_longitude']),
↳(x['end_station_latitude'], x['end_station_longitude'])),
axis=1)

```

```

[18]: ford_go_bike.sample(5)

```

```

[18]:      duration_sec      start_time      end_time \
6984          711 2019-02-28 09:24:34.441 2019-02-28 09:36:25.728
82017          638 2019-02-18 11:37:38.673 2019-02-18 11:48:17.377
102281         1285 2019-02-14 16:28:52.287 2019-02-14 16:50:17.537
119454          960 2019-02-11 17:42:49.634 2019-02-11 17:58:49.795
57360          338 2019-02-21 08:23:34.038 2019-02-21 08:29:12.072

      start_station_id      start_station_name \
6984          22.0          Howard St at Beale St
82017         243.0      Bancroft Way at College Ave
102281         16.0      Steuart St at Market St
119454         58.0      Market St at 10th St
57360         43.0  San Francisco Public Library (Grove St at Hyde...

      start_station_latitude  start_station_longitude  end_station_id \
6984          37.789756          -122.394643          28.0
82017          37.869360          -122.254337          171.0
102281          37.794130          -122.394430          87.0
119454          37.776619          -122.417385          126.0
57360          37.778768          -122.415929          350.0

      end_station_name  end_station_latitude \
6984  The Embarcadero at Bryant St          37.787168
82017  Rockridge BART Station          37.844279
102281  Folsom St at 13th St          37.769757
119454  Esprit Park          37.761634
57360  8th St at Brannan St          37.771431

      end_station_longitude  bike_id  user_type  member_birth_year \
6984          -122.388098      1142  Subscriber          1979

```



82017	-122.251900	4565	Subscriber	-9223372036854775808
102281	-122.415674	5331	Subscriber	1980
119454	-122.390648	3356	Subscriber	1988
57360	-122.405787	6094	Subscriber	1990

	member_gender	bike_share_for_all_trip	day	month	\
6984	Female	No	Thursday	February	
82017	NaN	No	Monday	February	
102281	Male	No	Thursday	February	
119454	Female	No	Monday	February	
57360	Male	No	Thursday	February	

	member_age	distance_km
6984	43	0.64
82017	-9223372036854773786	2.80
102281	42	3.29
119454	34	2.88
57360	32	1.21

```
[19]: #Let check for null values
      ford_go_bike.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 182380 entries, 0 to 183411
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   duration_sec                          182380 non-null  int64
1   start_time                            182380 non-null  datetime64[ns]
2   end_time                              182380 non-null  datetime64[ns]
3   start_station_id                      182183 non-null  float64
4   start_station_name                    182183 non-null  object
5   start_station_latitude                182380 non-null  float64
6   start_station_longitude               182380 non-null  float64
7   end_station_id                        182183 non-null  float64
8   end_station_name                      182183 non-null  object
9   end_station_latitude                 182380 non-null  float64
10  end_station_longitude                 182380 non-null  float64
11  bike_id                              182380 non-null  int64
12  user_type                             182380 non-null  object
13  member_birth_year                    182380 non-null  int64
14  member_gender                         174115 non-null  object
15  bike_share_for_all_trip               182380 non-null  object
16  day                                   182380 non-null  object
17  month                                 182380 non-null  object
18  member_age                           182380 non-null  int64
19  distance_km                           182380 non-null  float64
```

```
dtypes: datetime64[ns](2), float64(7), int64(4), object(7)
memory usage: 29.2+ MB
```

```
[20]: # Let check unique values in the member_gender
ford_go_bike['member_gender'].unique()
```

```
[20]: array(['Male', nan, 'Other', 'Female'], dtype=object)
```

```
[21]: # Lets confirm unique values in the day
ford_go_bike['day'].unique()
```

```
[21]: array(['Thursday', 'Wednesday', 'Tuesday', 'Monday', 'Sunday', 'Saturday',
          'Friday'], dtype=object)
```

```
[22]: ford_go_bike.dropna(inplace=True)
```

```
[23]: # Lets confirm that there are no empty "NaN" and other "None cells"
ford_go_bike.isna().sum()
```

```
[23]: duration_sec          0
start_time              0
end_time               0
start_station_id       0
start_station_name     0
start_station_latitude  0
start_station_longitude 0
end_station_id         0
end_station_name       0
end_station_latitude   0
end_station_longitude  0
bike_id                0
user_type              0
member_birth_year      0
member_gender          0
bike_share_for_all_trip 0
day                   0
month                 0
member_age             0
distance_km            0
dtype: int64
```

```
[24]: # Let check for outliers in the duration_sec column
ford_go_bike.duration_sec.describe([.99])
```

```
[24]: count    173920.000000
mean         704.223149
std         1645.440119
```

```

min          61.000000
50%          511.000000
99%          3176.000000
max          84548.000000
Name: duration_sec, dtype: float64

```

There seem to be outliers in the duration\_sec column. 99% for the duration\_sec fall under 3176 seconds. There is pretty awkward jump to 84548 as the maximum value suggests.

```

[25]: # Limiting duration in seconds to only 3200 seconds
ford_go_bike = ford_go_bike.query("duration_sec <= 3200")

```

```

[26]: #Test that duration in seconds is limited to only 3200 seconds
ford_go_bike.duration_sec.describe(percentiles=[.99])

```

```

[26]: count    172207.000000
      mean      613.170580
      std      426.521236
      min       61.000000
      50%      507.000000
      99%     2280.000000
      max     3200.000000
      Name: duration_sec, dtype: float64

```

### 1.3.1 Saving

```

[27]: ford_go_bike.to_csv('fordbike_clean.csv', index = False)
      fordbike_clean = pd.read_csv('fordbike_clean.csv');
      fordbike_clean.head(3)

```

```

[27]:
  duration_sec  start_time  end_time \
0          1585  2019-02-28 23:54:18.549  2019-03-01 00:20:44.074
1          1793  2019-02-28 23:49:58.632  2019-03-01 00:19:51.760
2          1147  2019-02-28 23:55:35.104  2019-03-01 00:14:42.588

  start_station_id  start_station_name  start_station_latitude \
0              7.0      Frank H Ogawa Plaza      37.804562
1             93.0  4th St at Mission Bay Blvd S      37.770407
2            300.0      Palm St at Willow St      37.317298

  start_station_longitude  end_station_id  end_station_name \
0          -122.271738      222.0  10th Ave at E 15th St
1          -122.391198      323.0  Broadway at Kearny
2          -121.884995      312.0  San Jose Diridon Station

  end_station_latitude  end_station_longitude  bike_id  user_type \

```

0	37.792714	-122.248780	4898	Subscriber
1	37.798014	-122.405950	5200	Subscriber
2	37.329732	-121.901782	3803	Subscriber

	member_birth_year	member_gender	bike_share_for_all_trip	day	\
0	1974	Male	Yes	Thursday	
1	1959	Male	No	Thursday	
2	1983	Female	No	Thursday	

	month	member_age	distance_km
0	February	48	2.41
1	February	63	3.33
2	February	39	2.03

[28]: *# high- leve overview of data shape and compositon*

```
print(fordbike_clean.shape)
print(fordbike_clean.dtypes)
print(fordbike_clean.head(10))
```

(172207, 20)

```
duration_sec      int64
start_time        object
end_time          object
start_station_id  float64
start_station_name object
start_station_latitude float64
start_station_longitude float64
end_station_id    float64
end_station_name  object
end_station_latitude float64
end_station_longitude float64
bike_id           int64
user_type         object
member_birth_year int64
member_gender     object
bike_share_for_all_trip object
day              object
month            object
member_age        int64
distance_km       float64
dtype: object
```

	duration_sec		start_time	end_time	\
0	1585	2019-02-28	23:54:18.549	2019-03-01 00:20:44.074	
1	1793	2019-02-28	23:49:58.632	2019-03-01 00:19:51.760	
2	1147	2019-02-28	23:55:35.104	2019-03-01 00:14:42.588	
3	1615	2019-02-28	23:41:06.766	2019-03-01 00:08:02.756	
4	1570	2019-02-28	23:41:48.790	2019-03-01 00:07:59.715	

5	1049	2019-02-28 23:49:47.699	2019-03-01 00:07:17.025
6	458	2019-02-28 23:57:57.211	2019-03-01 00:05:35.435
7	506	2019-02-28 23:56:55.540	2019-03-01 00:05:21.733
8	1176	2019-02-28 23:45:12.651	2019-03-01 00:04:49.184
9	395	2019-02-28 23:56:26.848	2019-03-01 00:03:01.947

	start_station_id	start_station_name \
0	7.0	Frank H Ogawa Plaza
1	93.0	4th St at Mission Bay Blvd S
2	300.0	Palm St at Willow St
3	10.0	Washington St at Kearny St
4	10.0	Washington St at Kearny St
5	19.0	Post St at Kearny St
6	370.0	Jones St at Post St
7	44.0	Civic Center/UN Plaza BART Station (Market St ...
8	127.0	Valencia St at 21st St
9	243.0	Bancroft Way at College Ave

	start_station_latitude	start_station_longitude	end_station_id \
0	37.804562	-122.271738	222.0
1	37.770407	-122.391198	323.0
2	37.317298	-121.884995	312.0
3	37.795393	-122.404770	127.0
4	37.795393	-122.404770	127.0
5	37.788975	-122.403452	121.0
6	37.787327	-122.413278	43.0
7	37.781074	-122.411738	343.0
8	37.756708	-122.421025	323.0
9	37.869360	-122.254337	252.0

	end_station_name	end_station_latitude \
0	10th Ave at E 15th St	37.792714
1	Broadway at Kearny	37.798014
2	San Jose Diridon Station	37.329732
3	Valencia St at 21st St	37.756708
4	Valencia St at 21st St	37.756708
5	Mission Playground	37.759210
6	San Francisco Public Library (Grove St at Hyde...	37.778768
7	Bryant St at 2nd St	37.783172
8	Broadway at Kearny	37.798014
9	Channing Way at Shattuck Ave	37.865847

	end_station_longitude	bike_id	user_type	member_birth_year \
0	-122.248780	4898	Subscriber	1974
1	-122.405950	5200	Subscriber	1959
2	-121.901782	3803	Subscriber	1983
3	-122.421025	6329	Subscriber	1989
4	-122.421025	6548	Subscriber	1988

5	-122.421339	6488	Subscriber	1992
6	-122.415929	5318	Subscriber	1996
7	-122.393572	5848	Subscriber	1993
8	-122.405950	5328	Customer	1990
9	-122.267443	4786	Subscriber	1988

	member_gender	bike_share_for_all_trip	day	month	member_age \
0	Male	Yes	Thursday	February	48
1	Male	No	Thursday	February	63
2	Female	No	Thursday	February	39
3	Male	No	Thursday	February	33
4	Other	No	Thursday	February	34
5	Male	No	Thursday	February	30
6	Female	Yes	Thursday	February	26
7	Male	No	Thursday	February	29
8	Male	No	Thursday	February	32
9	Male	No	Thursday	February	34

	distance_km
0	2.41
1	3.33
2	2.03
3	4.53
4	4.53
5	3.66
6	0.98
7	1.61
8	4.78
9	1.21

### 1.3.2 What is the structure of your dataset?

After preliminary wrangling and feature engineering we are now left with (172207 bike rides that occurred in the year 2019 with 20 features(duration\_sec, start\_time, end\_time, start\_station\_id, start\_station\_name, start\_station\_latitude, start\_station\_longitude, end\_station\_id, end\_station\_name, end\_station\_latitude, end\_station\_longitude, bike\_id, user\_type, member\_birth\_year, member\_gender, bike\_share\_for\_all\_trip, day, month, member\_age, distance\_km). 5 ints (duration\_sec, member\_age, member\_birth\_year, and bike\_id )

7 floats(distance\_km, start\_station\_id, start\_station\_latitude, start\_station\_longitude, end\_station\_id, end\_station\_latitude, end\_station\_longitude)

9 string objects ( user\_type, end\_station\_name, start\_time, end\_time, start\_station\_name, member\_gender, bike\_share\_for\_all\_trip, day, month)

The feature engineered columns:

- day

- month
- member\_age
- distance\_km

### 1.3.3 What is/are the main feature(s) of interest in your dataset?

I'm most interested in determining which features are most effective at predicting the duration of trips in the dataset.

### 1.3.4 What features in the dataset do you think will help support your investigation into your feature(s) of interest?


I expect that day, month, member\_age, distance\_km would affect ride duration\_sec.

## 1.4 Univariate Exploration

We will begin by examining the distribution of the main variable of interest: duration\_sec

### 1.4.1 In what duration(sec/mins) do we have the most number of trips ?

```
[29]: # lets start with a standard-scaled plot
binsize = 500
bins = np.arange(0, ford_go_bike['duration_sec'].max()+binsize, binsize )
plt.figure(figsize=[8, 5])
plt.hist(data = ford_go_bike, x = 'duration_sec', bins=bins);
plt.title('Distribution of Bike ride trip durations')
plt.xlabel('Duration_sec')
plt.ylabel('No of Trips')
plt.axis([-500, 10000, 0, 90000])
plt.show()
```




```
[30]: # There is a long tail in hte distribution , so let us put it on a log scale
log_binsize = 0.05
bins = 10 ** np.arange(2.4, np.log10(ford_go_bike['duration_sec'].max()) +
    ↳log_binsize, log_binsize)

plt.figure(figsize=[8, 5])
plt.hist(data = ford_go_bike, x = 'duration_sec', bins = bins)
plt.title('Distribution of Bike ride trip durations')
plt.xlabel('Duration_sec')
plt.ylabel('No of Trips')
plt.xscale('log')
plt.xticks([500, 1e3, 2e3, 5e3, 1e4], ['8.3mins', '16.6mins', '33.3mins', '1hr.
    ↳23mins', '2hr.47mins'])
plt.axis([0, 10000, 0, 15000])
plt.show()
```

```
/tmp/ipykernel_16761/1134385293.py:12: UserWarning: Attempted to set non-
positive left xlim on a log-scaled axis.
Invalid limit will be ignored.
    plt.axis([0, 10000, 0, 15000])
```





Initially `duration_sec` had a long-tailed distribution with very few rides at the high end duration. When plotted on a log-scale, most trip duration occurred in less than 33mins 20 secs (2000 seconds ). The highest record of trips occurred in 10mins (600 secs). The number of trips initially increases from around 8000 values at 0 to 12000 values at around 600 seconds, but then begins to fall, reaching below 2000 values in less than 33 minutes.

#### 1.4.2 What day of the week have the most bike rides ?

Let us visualize day a categorical variable using Bar plt

```
[31]: #Let define a function that would be used regularly to label our visualizations
def label_visual(x_label, y_label, title):
    plt.xlabel(x_label)
    plt.ylabel(y_label)
    plt.title(title)
```

```
[32]: # bar plot of `day` categorical variable


#let order the chat base on weekdays
weekdays = ["Sunday", "Monday", "Tuesday", "Wednesday", "Thursday", "Friday",
↪ "Saturday"]

# plotting with one base color
```

```
base_color = sb.color_palette()[0]

sb.countplot(data = fordbike_clean, x = 'day', order=weekdays,
             color=base_color);
label_visual("Day of the Week", "No of Trips", "Bike usage distribution by",
             Weekday")

plt.xticks(rotation=90);
```



Thursdays had the highest usage, followed by Tuesdays and Fridays. Saturdays and Sundays saw a significant decrease in usage. This suggests that bikes are mostly used during the week and/or that people prefer to stay at home on weekends, resulting in less use of bikes.

### 1.4.3 What is the age distribution of the of the bike riders ?

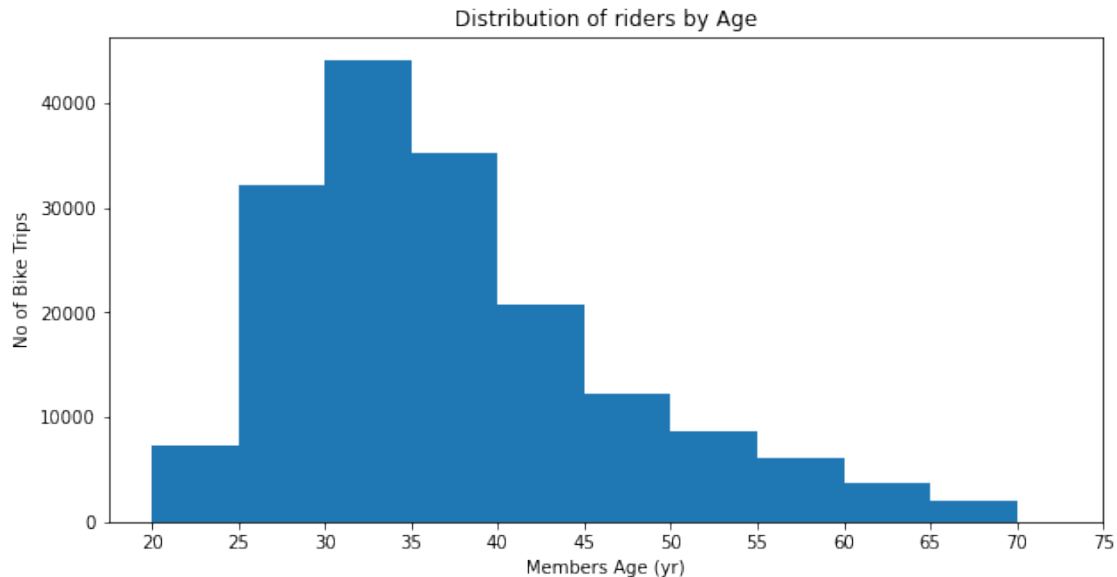
Histogram of the discrete variable age

```
[33]: plt.figure(figsize = [10,5])
      bin_age = np.arange(20, fordbike_clean.member_age.max()+5, 5)
```

```

ticks = [20, 25, 30, 35, 40, 45, 50, 55, 60, 65, 70, 75]
labels = ['{}'.format(i) for i in ticks]
plt.hist(data = fordbike_clean, x= 'member_age', bins= bin_age);
plt.xticks(ticks, labels);
label_visual('Members Age (yr)', 'No of Bike Trips', 'Distribution of riders by Age')

```



The histogram shows that most bike riders fall between the ages of 25 to 40 year. Individual between 30-35 have the heighest number of rides.

#### 1.4.4 What is the distribution of users type ?

[34]: *# Pie plot distribution of the user type (nominal categorical variable)*

```

fig, ax = plt.subplots(figsize=(10, 5))

user_count = fordbike_clean.user_type.value_counts()


patches, texts, pcts = ax.pie(
    user_count, labels=user_count.index, autopct='%.1f%%',
    wedgeprops={'linewidth': 3.0, 'edgecolor': 'white'},
    textprops={'size': 'large'},
    startangle=90)
# For each wedge, set the corresponding text label color to the wedge's
# face color.
for i, patch in enumerate(patches):

```

```

    texts[i].set_color(patch.get_facecolor())
plt.setp(pcts, color='white')
plt.setp(texts, fontweight=100)
ax.set_title('Distribution for User types')
plt.tight_layout()

```



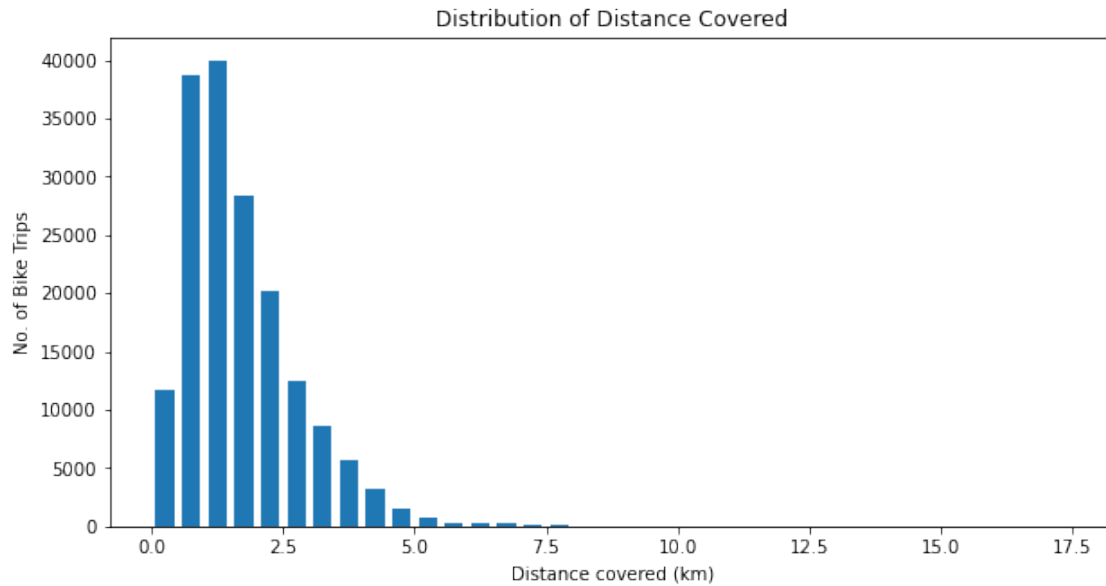
Subscriber users account for 90.9% of all bike rides, while customer users account for only 9.1% of all bike rides.

#### 1.4.5 What is the most distance covered by number of bike trips ?

```

[35]: plt.figure(figsize=[10, 5])
      bins_distance = np.arange(0, fordbike_clean.distance_km.max()+2, 0.5)
      plt.hist(data = fordbike_clean, x = 'distance_km', bins = bins_distance,
               ↪rwidth=0.7)
      label_visual('Distance covered (km)', 'No. of Bike Trips', 'Distribution of_
               ↪Distance Covered')

```



The histogram above shows the distribution of distance covered (km). Majority of the distances covered ranged from 0.1 to 2.5 km. This demonstrates that the majority of trips were not long distances.

#### 1.4.6 Which gender has the most bike rides ?

```
[36]: #Which gender has the most bike rides ?
print("%3.1f"%(((fordbike_clean['member_gender'].value_counts()[0])/
↳fordbike_clean['member_gender'].value_counts().sum()*100))
print("%3.1f"%(((fordbike_clean['member_gender'].value_counts()[1])/
↳fordbike_clean['member_gender'].value_counts().sum()*100))
print("%3.1f"%(((fordbike_clean['member_gender'].value_counts()[2])/
↳fordbike_clean['member_gender'].value_counts().sum()*100))
```


74.6

23.4

2.0

```
[37]: #pie plot of member_gender distribution
plt.figure(figsize=[10, 5])
sorted_counts = fordbike_clean['member_gender'].value_counts()

plt.pie(sorted_counts, labels = sorted_counts.index, startangle = 90,
↳counterclock = False, wedgeprops={'edgecolor': 'white'}, textprops={'size':
↳'large'});
plt.axis('square');
```



The plot shows that we have mostly Male riders constituting 75%, Female riders 23% and others 2%

**1.4.7 Discuss the distribution(s) of your variable(s) of interest. Were there any unusual points? Did you need to perform any transformations?**

Initially, after visual and programmatic assesement we made serious attempts to eliminate outliers in some columns such as `duration_sec` , `member_age` and a few other columns.

**1.4.8 Of the features you investigated, were there any unusual distributions? Did you perform any operations on the data to tidy, adjust, or change the form of the data? If so, why did you do this?**

We limited most of the dataset to the values of 99% of the entire dataset. This was done to eliminate outliers in the dataset.

## 1.5 Bivariate Exploration

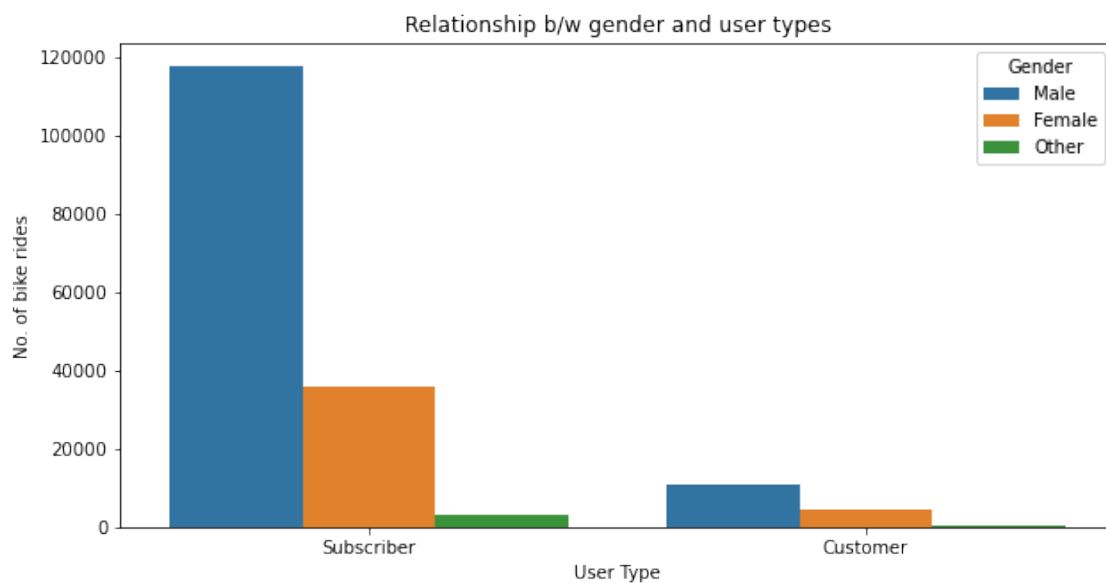
In this section, we shall investigate relationships between different pairs of variables in our dataset

### 1.5.1 What is the relationship between gender and user type

```
[38]: # We use a cluster bar chart to show the relationship btw to qualitative
      ↪ variables.
      plt.figure(figsize=[10, 5])

      sb.countplot(data =fordbike_clean, x = 'user_type', hue = 'member_gender')
      legend = plt.legend([ "Male","Female", "Other"])
      legend.set_title("Gender")

      label_visual("User Type", "No. of bike rides", "Relationship b/w gender and
      ↪ user types")
```




The plot demonstrates that there are more male subscribers than male customers, while there are little or no others gender in customer users and a few male and female customers

### 1.5.2 Which user type spends the most time riding ?

We shall look at the relationship between duration\_sec (quantitative) and user\_type(qualitative) using a Violin Plot.

```
[39]: #Violent plot showing the relationship btw duration_sec and user_type
      # set and create the figure size
      plt.figure(figsize=[10, 5])
      users = ["Customer", "Subscriber"]
      sb.violinplot(data=fordbike_clean, x = "user_type", y="duration_sec",
      ↪ order=users)
```

```
label_visual("User Types", "Duration in seconds", "Relationship btw duration_sec_
↳and user_type")
```




The figure demonstrates the relationship btw duration\_sec and user\_type. More Customers types send greater time between 500sec and 1000secs. While Subscriber users mostly spend an average of 500sec.

### 1.5.3 In which day of the Week do riders spends the most time.?

We shall look at the relationship between duration\_sec (quantitative) and weekday (qualitative) using a Bar Plot.

```
[40]: #relationship between duration_sec (quantitative) and weekday (qualitative)
↳using a bar Plot.
plt.figure(figsize = [10, 5])
weekdays = ["Sunday", "Monday", "Tuesday", "Wednesday", "Thursday", "Friday",
↳"Saturday"]
sb.barplot(data= fordbike_clean, x = "day", y="duration_sec" , order=weekdays,
↳color=base_color)
label_visual("Days of the Week", "Duration (sec)", "Relationship b/w
↳Duration_sec and day")
```



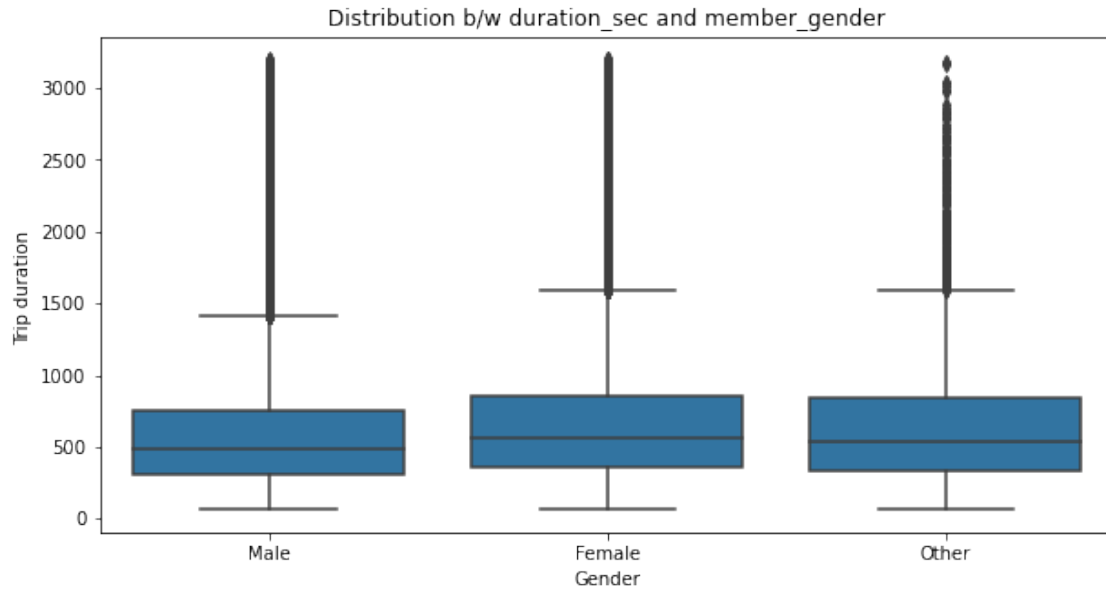


The graph shows the relationship between duration\_sec (quantitative) and weekday (qualitative) using a Bar Plot. Riders spend more time riding on Sunday, Saturday than on workdays (Mon - Friday)

#### 1.5.4 Which gender spends the most time riding bikes ?

Using a Box Plot we shall examine the distribution between trip duration (quantitative variable) and gender (qualitative variable)

```
[41]: #Box plot showing distribution b/w duration_sec and member_gender
plt.figure(figsize = [10, 5])
sb.boxplot(data=fordbike_clean, x='member_gender', y= 'duration_sec',
           color=base_color)
label_visual('Gender', 'Trip duration ', 'Distribution b/w duration_sec and',
           member_gender)
```




From the distribution, Females spent more time on their bikes than men on average.

### 1.5.5 What is the relationship between distance covered and trip duration. ?

We use Seaborn's `regplot()` function that combines scatterplot creation with regression function fitting:

```
[42]: #the relationship between distance covered and trip duration
plt.figure(figsize= [20, 10])
sb.regplot(data = fordbike_clean, x = 'distance_km', y='duration_sec')
label_visual('Distance covered (km)', ' Duration (sec)', 'relationship between_
↳distance covered and trip duration')
```



As expected, the regression line in the scatter plot shows a positive correlation between distance covered and duration

#### 1.5.6 Talk about some of the relationships you observed in this part of the investigation. How did the feature(s) of interest vary with other features in the dataset?

There are more male subscribers than male customers, while there are little or no others gender in customer users and a few male and female customers. More Customers types send greater time between 500sec and 1000secs. While Subscriber users mostly spend an average of 500sec. Riders spend more time riding on Sunday, Saturday than on workdays(Mon - Friday). As expected there is a positive correlation between distance covered and duration.

#### 1.5.7 Did you observe any interesting relationships between the other features (not the main feature(s) of interest)?

It is interesting to know that Riders spend more time riding on Sunday, Saturday than on workdays(Mon - Friday)

[ ]:


### 1.6 Multivariate Exploration

In this section we shall investigate the relationship in more than 2 variables at once.

### 1.6.1 What is the relationship between duration, distance and age ?

```
[43]: plt.figure(figsize= [20, 20])
cat_markers = [['Male', "^"],
               ['Female', "v"],
               ['Other', "o"]]

for cat, marker in cat_markers:
    df_cat = fordbike_clean[fordbike_clean['member_gender'] == cat]
    plt.scatter(data = df_cat, x = 'duration_sec', y = 'distance_km', marker =_
    ↪marker)
plt.legend(['Male', 'Female', 'Other'])
label_visual('duration_sec', 'distance_km', 'Relationship between duration,_
    ↪distance and gender')
```

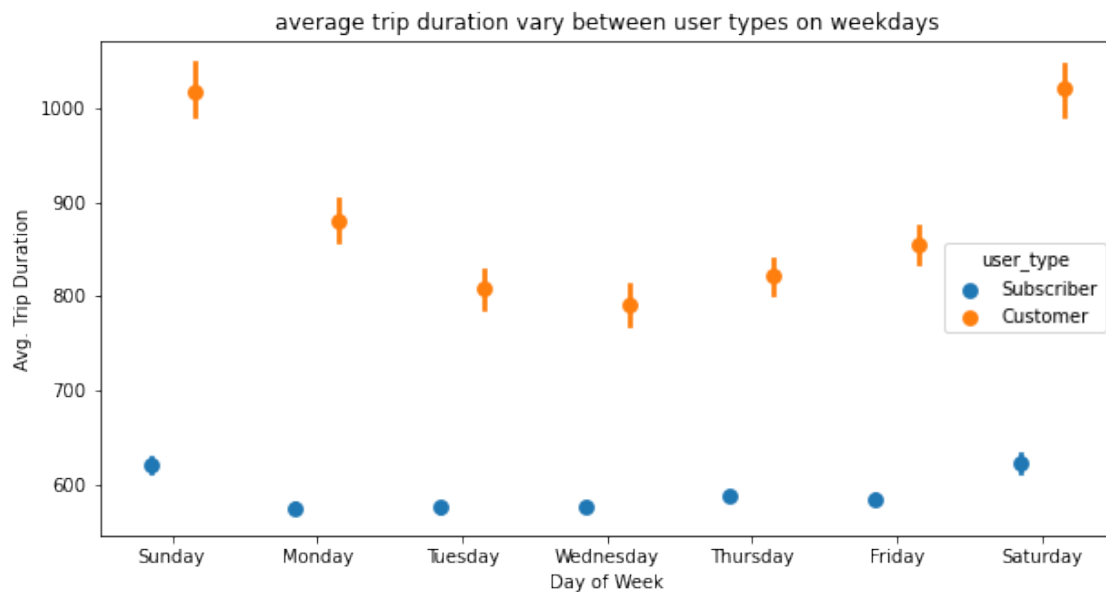


The graph shows the relationship between duration, distance and gender. There are more males covering longer distances than females. There are more females spending more time riding bikes.

### 1.6.2 What is the do the average trip duration vary between user types on weekdays?

```
[44]: # Let look at the relationship between user_type, duration_sec and day.
plt.figure(figsize= [10, 5])


sb.pointplot(data=fordbike_clean, x='day', y='duration_sec',
             hue='user_type', dodge=0.3, linestyle="", order=weekdays);
label_visual('Day of Week', 'Avg. Trip Duration', 'average trip duration vary_
↪between user types on weekdays')
```



Generally more Customers spend more time riding bikes than subscribers. Both customers and subscribers spend more time on Sundays and Saturdays. There is not much preference in work-days (Mon - Fri) for subscribers however, for customers most trip duration occurred on Monday.

Finally let examine the Correlation matrices among all numerical variables

```
[45]: plt.figure(figsize=[15, 8])
sb.heatmap(fordbike_clean.corr(), annot= True, fmt='.2f',
           cmap = 'vlag_r', center = 0);
```



There is a strong relationship between duration in seconds and minutes and distance traveled in kilometers.

### 1.6.3 Talk about some of the relationships you observed in this part of the investigation. Were there features that strengthened each other in terms of looking at your feature(s) of interest?

Males are more likely to travel longer distances than females. Females are riding their bikes more frequently. Customers ride their bikes more than subscribers. On Sundays and Saturdays, both customers and subscribers spend more time. Subscribers have little preference for workdays (Mon-Fri), but customers have the most trip duration on Monday. The duration in seconds and minutes and the distance traveled in kilometers have a strong relationship.

### 1.6.4 Were there any interesting or surprising interactions between features?

Males are more likely to travel longer distances than females while females ride their bikes more frequently than males.

[ ]:

## 1.7 Conclusions

In conclusion, bike riders tend to be between the ages of 25 and 40, who use their bikes the most frequently and ride for the longest periods of time during the workweek. The trend among riders is to travel exclusively by bicycle on Thursdays. Males are more likely to travel longer distances than females while females ride their bikes more frequently than males.

A lot of effort was put into wrangling to ensure the best outcome of the analysis.

[ ]: